

# Parallel Computing Workshop

■ Ms. Kranti Ingale (CS22D003)  
Prime Minister's Research Fellows (**PMRF**)  
PACE Lab - CSE, IIT Madras

# Roadmap

- Introduction to GPU
- CUDA Program Flow and CPU-GPU Communication
- Thread organization (Grids, Blocks, Threads, 1D/2D)
- CUDA Memory Model
- CUDA Functions
- CUDA Thrust

# Roadmap

- Introduction to GPU
- CUDA Program Flow and CPU-GPU Communication
- Thread organization (Grids, Blocks, Threads, 1D/2D)
- CUDA Memory Model
- CUDA Functions
- CUDA Thrust

# Introduction

- A Graphics Processing Unit (GPU) is a microprocessor designed specifically for the processing of 3D graphics.
- GPU forms the heart of modern graphics card, relieving the CPU (Central Processing Unit) of much of the graphics processing load.
- GPU has become integral part of today's computer systems.
- Over the last few years, there has been remarkable increase in performance and capabilities of GPU



## Why GPU?

- To provide separate graphics processing resources
- To relieve some of the burden of the main system resources



# GPU VS CPU

## GPU

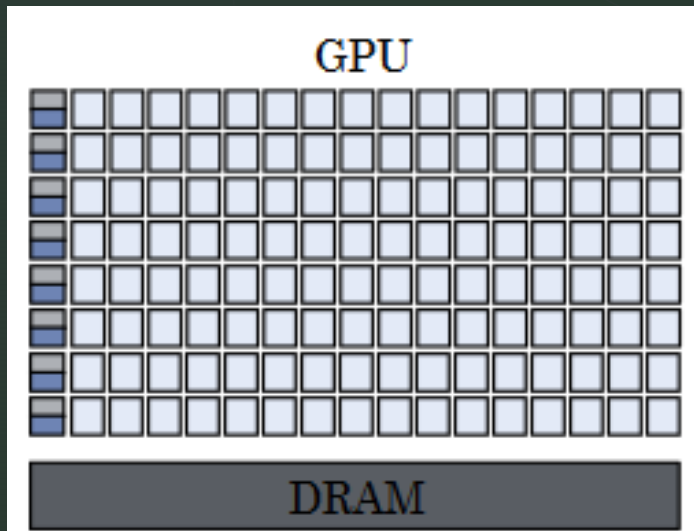
- Tailored for highly parallel operations
- GPUs have many parallel execution units
- Many-core

## CPU

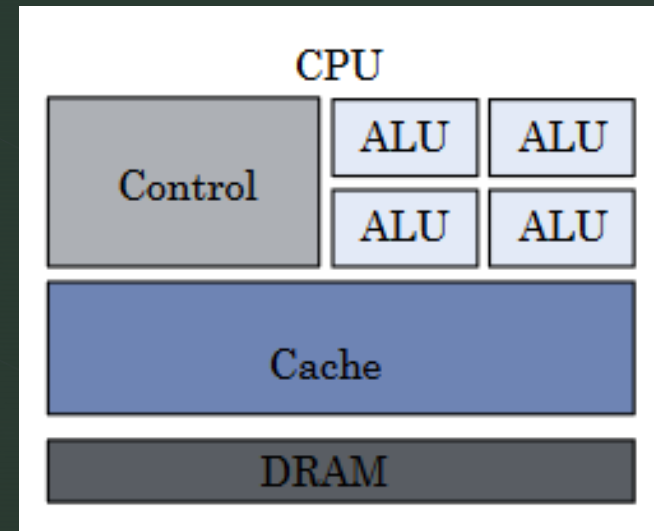
- CPU executes programs sequentially
- CPUs have few execution unit
- Multi-core

# GPU VS CPU

GPU



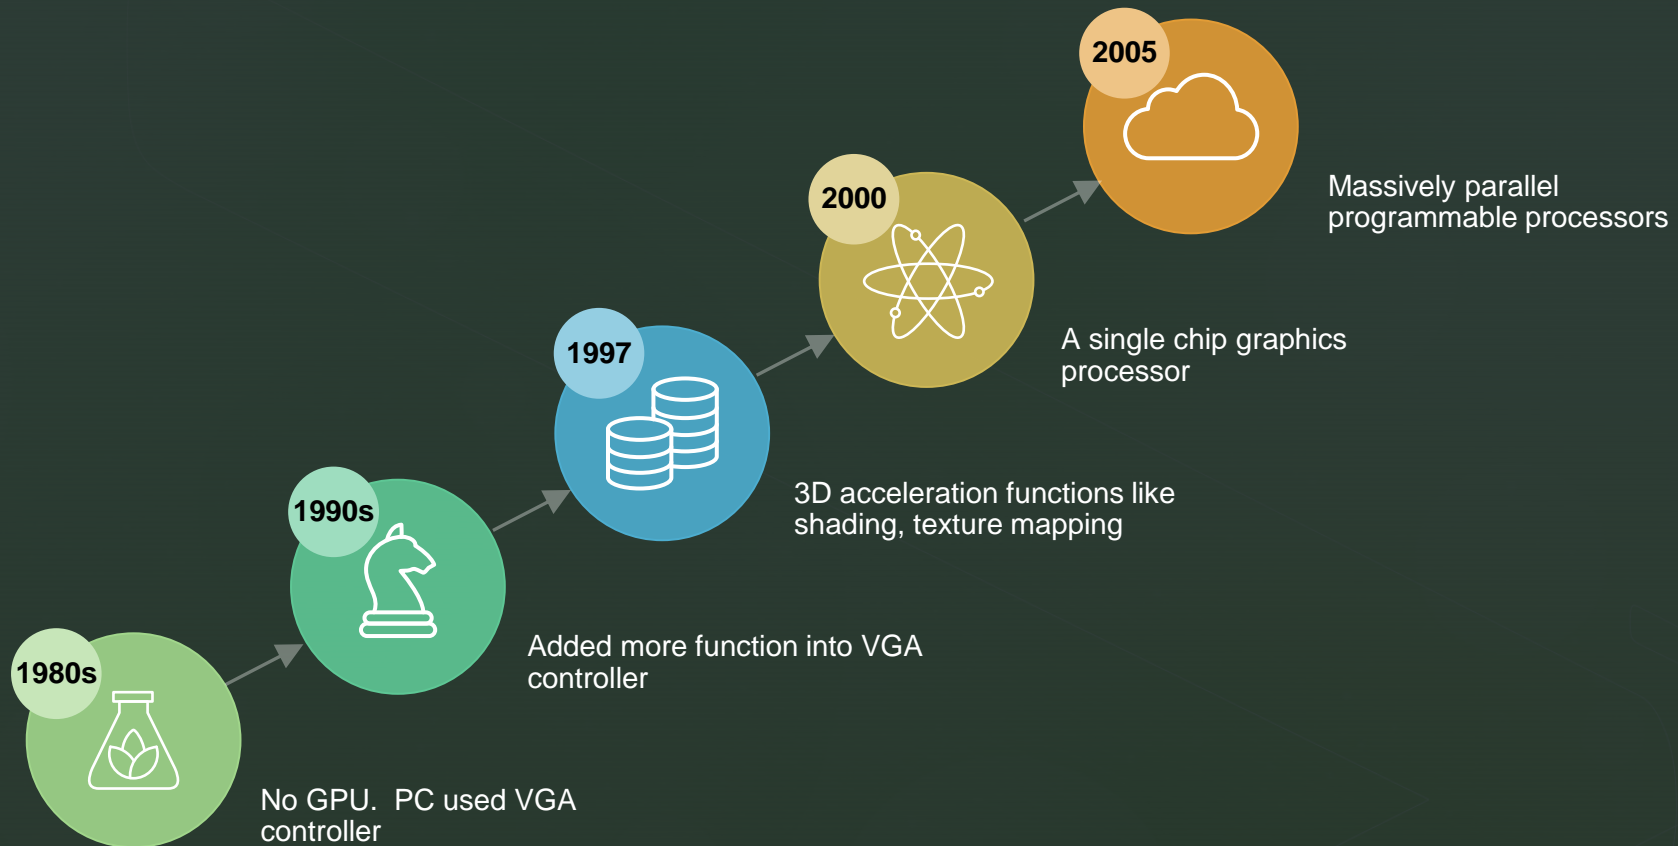
CPU





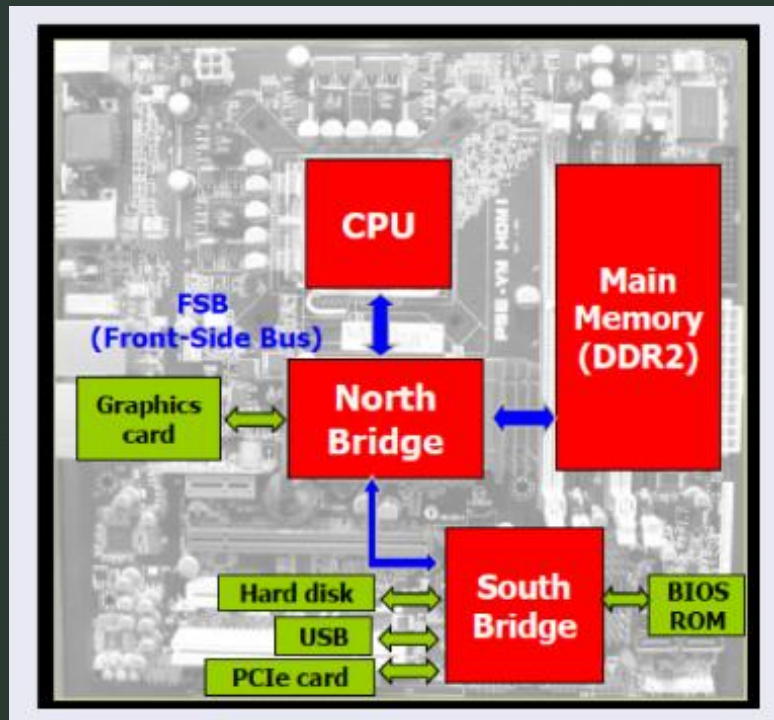
# GPU Evolution

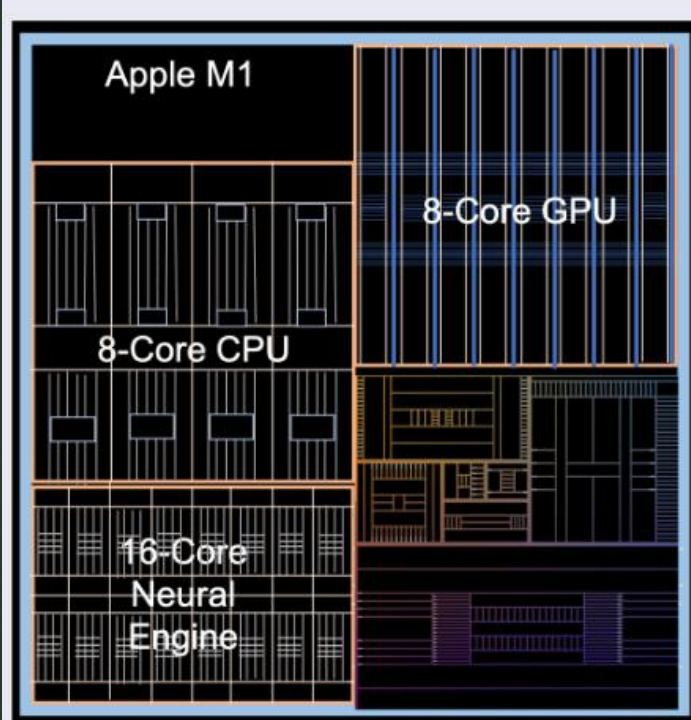
Highly parallel, highly multithreaded multiprocessor optimized for graphic computing and other applications





## GPUs: Discrete to Integrated





# GPU Vendors

- NVIDIA
- AMD
- Intel
- Qualcomm
- ARM
- ...

# GPU Programming

- GPU Programming empowers developers to harness the immense computational power of GPUs
- General Purpose computation (GP GPU): using GPU in applications other than 3D graphics
- Applications
  - Game effects (FX) physics, image processing
  - Physical modeling, computational engineering, matrix algebra, convolution, correlation, sorting

# GPU Languages

- CUDA (compute unified device language)
- OpenCL (open computing language)
- OpenACC (open accelerator)
- Sycl

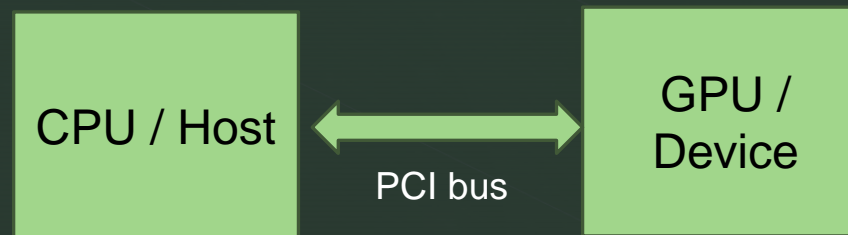
# GPU Languages

- CUDA (compute unified device language)
- OpenCL (open computing language)
- OpenACC (open accelerator)
- Sycl



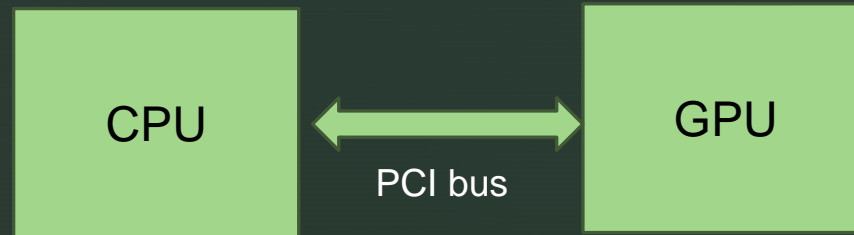
## Terminology

- **Host** : The CPU and its memory (host memory)
- **Device** : The GPU and its memory (device memory)



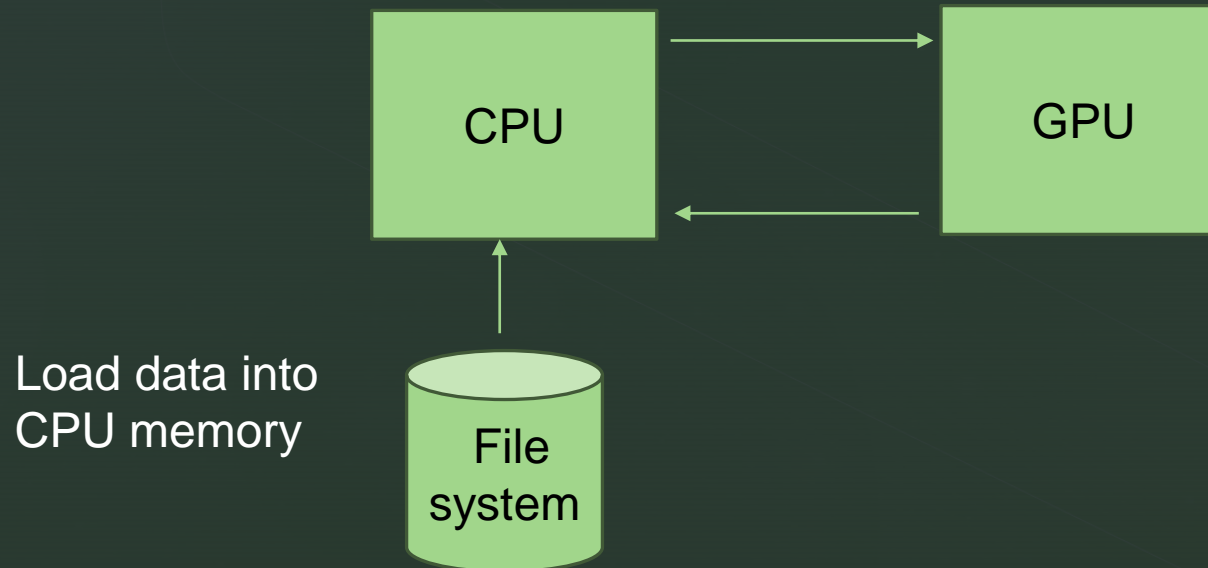


## Separate Memories

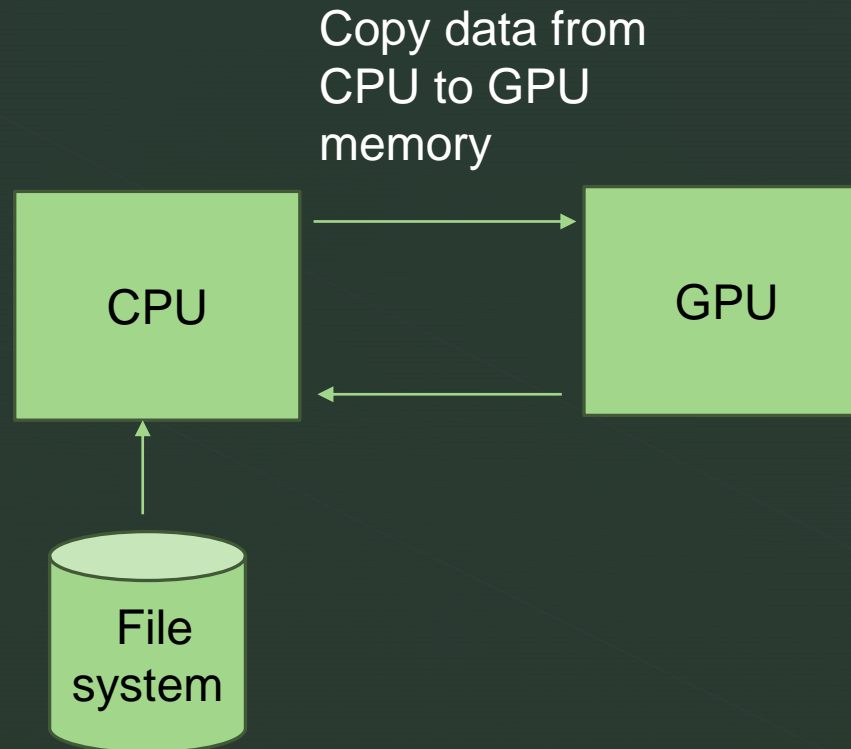


- The CPU and its corresponding discrete GPUs possess distinct physical memory (RAM).
- Direct access to a variable in CPU memory is not possible from a GPU .
- A programmer is responsible for maintaining copies of variables, and it is their duty to ensure that these copies remain in sync.

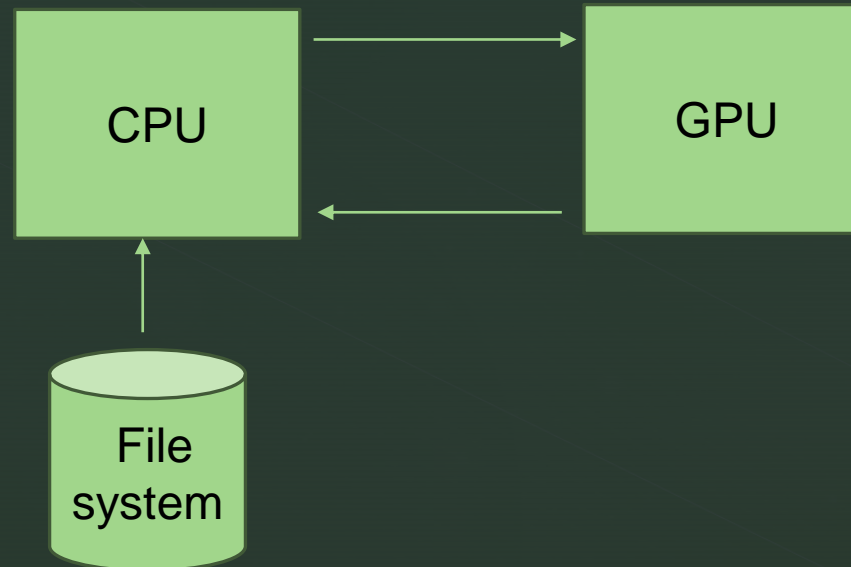
## Programming using CUDA



# Programming using CUDA



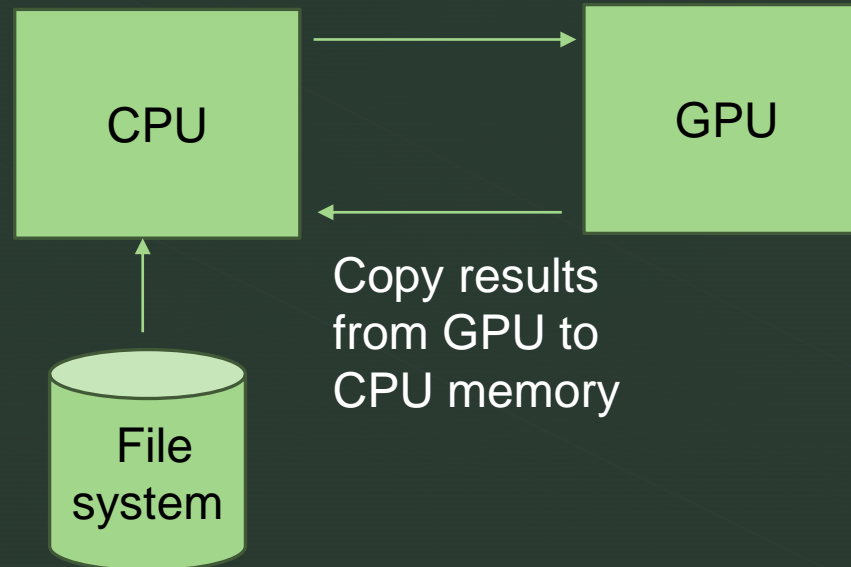
# Programming using CUDA



Execute GPU  
kernel

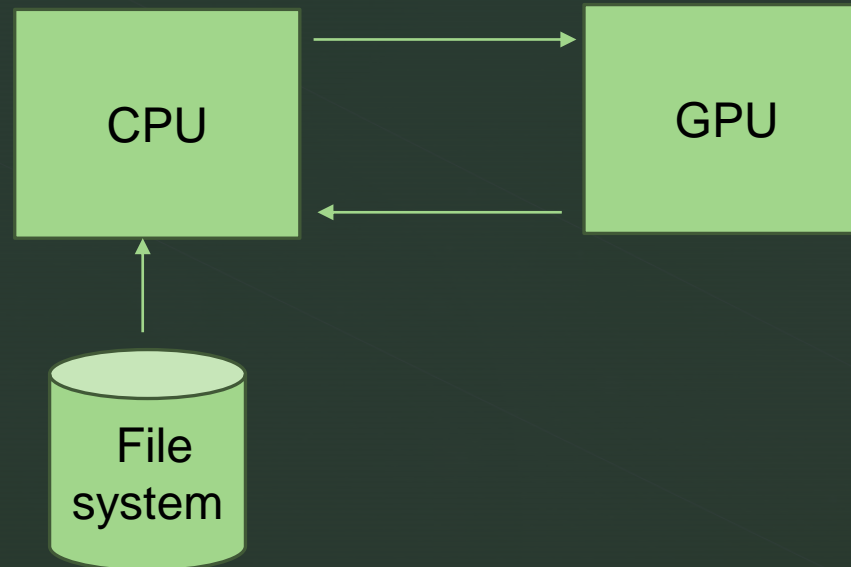


## Programming using CUDA

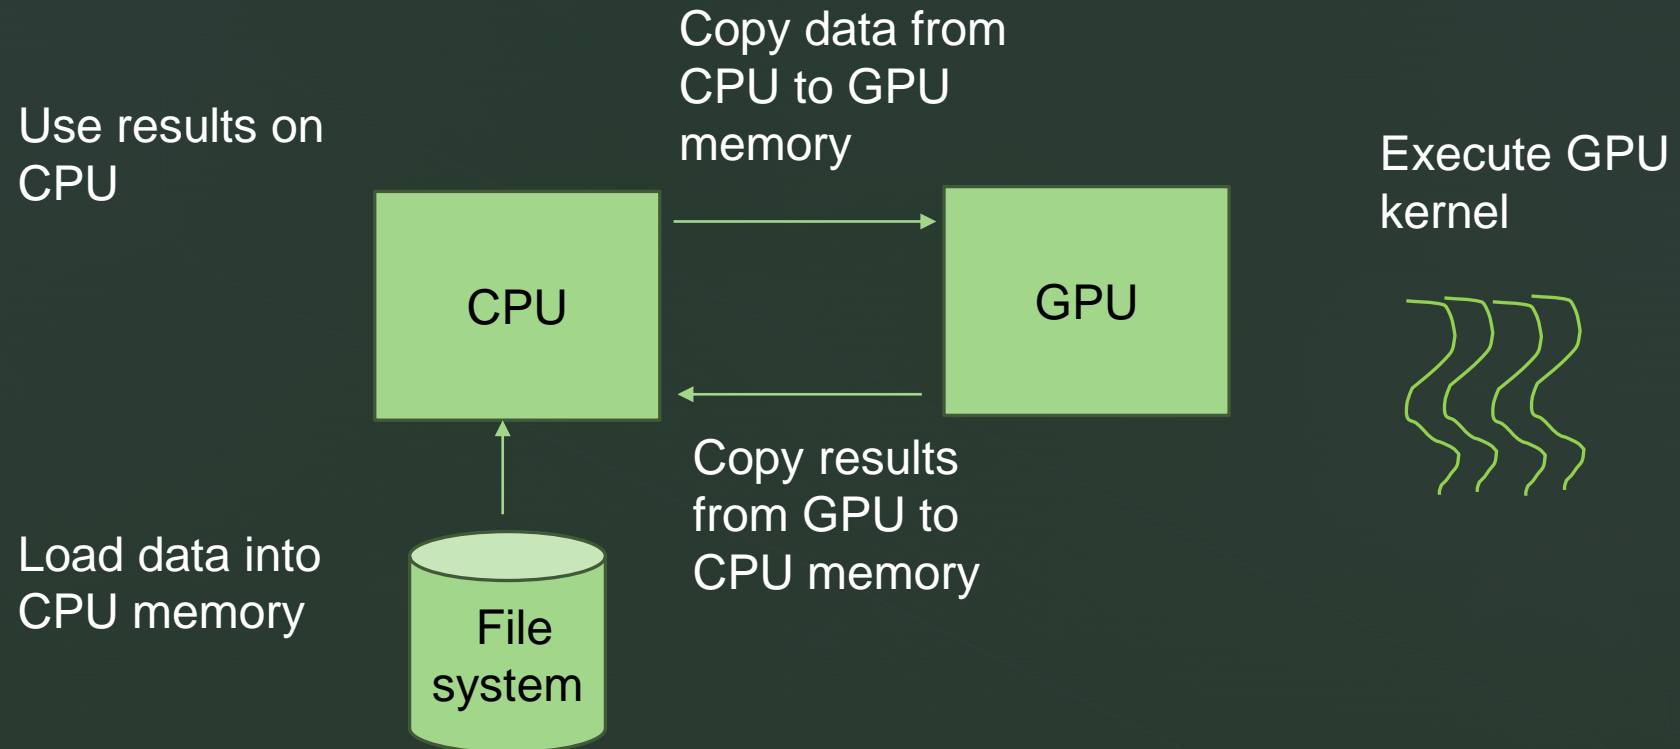


# Programming using CUDA

Use results on  
CPU



# Programming using CUDA



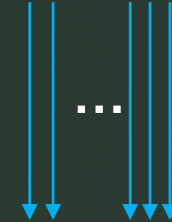


# Programming using CUDA

- Load data into CPU memory
- Copy data from CPU to GPU memory
- Call GPU kernel
- Copy results from GPU to CPU memory
- Use results on CPU

# Programming using CUDA

- `do_something_on_host();`
- `kernel<<<nBlk, nTid>>>(args);`
- `cudaDeviceSynchronize();`
- `do_something_else_on_host();`



Highly parallel





Thank You

