

Parallel Computing Workshop

Ms. Kranti Ingale

Prime Minister's Research Fellows (**PMRF**) Scholar

PACE Lab - CSE, IIT Madras

Roadmap

- Introduction to GPU
- CUDA Program Flow and CPU-GPU Communication
- Thread organization (Grids, Blocks, Threads, 1D/2D)
- CUDA Memory Model
- CUDA Functions
- CUDA Thrust

Roadmap

- Introduction to GPU
- CUDA Program Flow and CPU-GPU Communication
- Thread organization (Grids, Blocks, Threads, 1D/2D)
- CUDA Memory Model
- CUDA Functions
- CUDA Thrust

Hello World

```
#include <stdio.h>

int main() {
    printf("Hello World.\n");
    return 0;
}
```

Hello World

```
#include <stdio.h>
int main() {
printf("Hello World.\n");
return 0;
}
```

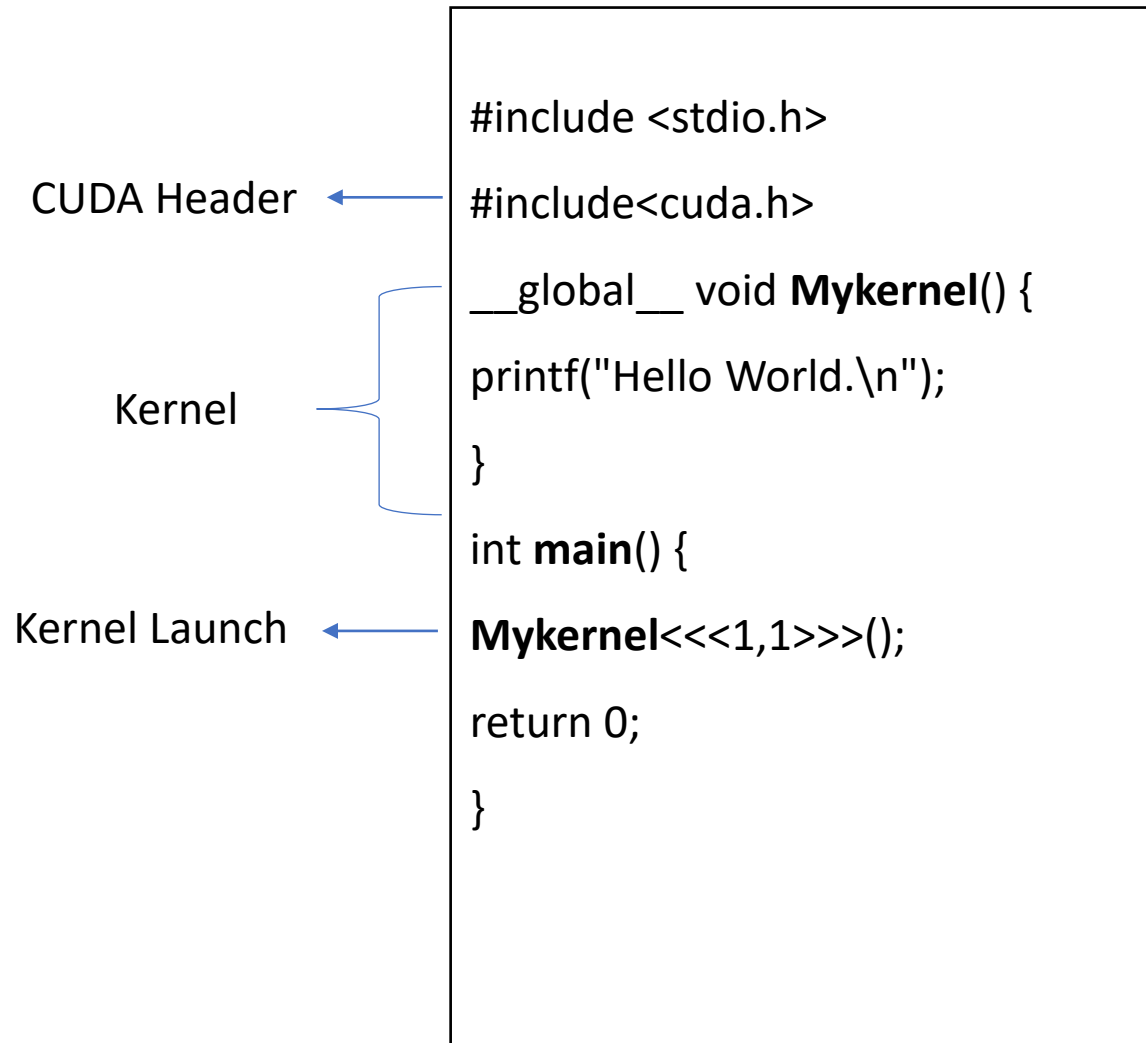
Compile: nvcc <ProgramName>.cu

Run: a.out

GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
__global__ void Mykernel() {
    printf("Hello World.\n");
}
int main() {
    Mykernel<<<1,1>>>();
    return 0;
}
```

GPU Hello World program



GPU Hello World program

CUDA Header



```
#include <stdio.h>
#include<cuda.h>
__global__ void Mykernel() {
printf("Hello World.\n");
}

int main() {
Mykernel<<<1,1>>>();
return 0;
}
```

Note :

cuda.h is header file needs to include in order to run CUDA programs

GPU Hello World program

Kernel

```
#include <stdio.h>
#include <cuda.h>

__global__ void Mykernel() {
    printf("Hello World.\n");
}

int main() {
    Mykernel<<<1,1>>>();
    return 0;
}
```

Note :

Kernel is Function that :

- Runs on the device
- Is called from host code

GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
__global__ void Mykernel() {
    printf("Hello World.\n");
}
int main() {
    Mykernel<<<1,1>>>();
    return 0;
}
```

Kernel Launch



Note :

Kernel are executed N times in parallel by N different *CUDA threads*.

GPU Hello World program

CUDA Header ← `#include <stdio.h>`
`#include <cuda.h>`

Kernel {
 `__global__ void Mykernel() {`
 `printf("Hello World.\n");`
 `}`

Kernel Launch ← `int main() {`
 `Mykernel<<<1,1>>>();`
 `return 0;`
}

Output :

Compile: `nvcc hello.cu`

Run: `./a.out`

----No Output----

GPU Hello World program

```
#include <stdio.h>
#include<cuda.h>

__global__ void Mykernel() {
printf("Hello World\n");
}

int main() {
Mykernel<<<1,1>>>();
cudaDeviceSynchronize();
return 0;
}
```

Output:

Hello World

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>

__global__ void Mykernel() {
    printf("Hello World\n");
}

int main() {
    Mykernel<<<1,1>>>();
    Mykernel<<<1,1>>>();
    cudaDeviceSynchronize();
    return 0; }
```

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
__global__ void Mykernel() {
    printf("Hello World\n");
}
int main() {
    Mykernel<<<1,1>>>();
    Mykernel<<<1,1>>>();
    cudaDeviceSynchronize();
    return 0; }
```

Output:

Hello World
Hello World

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>

__global__ void Mykernel() {
    printf("Hello World\n");
}

int main() {
    Mykernel<<<1,1>>>();
    cudaDeviceSynchronize();
    Mykernel<<<1,1>>>();
    return 0; }
```

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>

__global__ void Mykernel() {
    printf("Hello World\n");
}

int main() {
    Mykernel<<<1,1>>>();
    cudaDeviceSynchronize();
    Mykernel<<<1,1>>>();
    return 0; }
```

Output:

Hello World

Hello World (Optional)

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    cudaDeviceSynchronize();  
    return 0; }
```

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    cudaDeviceSynchronize();  
    return 0; }
```

Output:

Hello World 1
Hello World 2

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    cudaDeviceSynchronize();  
    Printf("Inside main\n");  
    return 0; }
```

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    cudaDeviceSynchronize();  
    Printf("Inside main\n");  
    return 0; }
```

Output:

Hello World 1

Hello World 2

Inside main

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    Printf("Inside main\n");  
    cudaDeviceSynchronize();  
    return 0; }
```

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    Printf("Inside main\n");  
    cudaDeviceSynchronize();  
    return 0; }
```

Output:

Hello World 1

Hello World 2

Inside main

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    Printf("Inside main\n");  
    cudaDeviceSynchronize();  
    return 0; }
```

Output:

```
Hello World 1  
Hello World 2  
Inside main  
  
Hello World 1  
Inside main  
Hello World 2
```

More GPU Hello World program

```
__global__ void Mykernel1() {  
    printf("Hello World 1\n");  
}  
  
__global__ void Mykernel2() {  
    printf("Hello World 2\n");  
}  
  
int main() {  
    Mykernel1<<<1,1>>>();  
    Mykernel2<<<1,1>>>();  
    Printf("Inside main\n");  
    cudaDeviceSynchronize();  
    return 0; }
```

Output:

```
Hello World 1  
Hello World 2  
Inside main  
  
Hello World 1  
Inside main  
Hello World 2  
  
Inside main  
Hello World 1  
Hello World 2
```


More GPU Hello World program

```
__global__ void Mykernel() {  
    printf("Hello World \n");  
}  
  
int main() {  
    Mykernel<<<1,1>>>();  
    Printf("main one\n");  
    Mykernel<<<1,1>>>();  
    Printf("main two\n");  
    cudaDeviceSynchronize();  
    Printf("main three\n");  
    return 0; }
```

Identify which prints execute in parallel

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
__global__ void Mykernel() {
    printf("Hello World\n");
}
int main() {
    Mykernel<<<1,32>>>();
    return 0;
}
```

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>

__global__ void Mykernel() {
    printf("Hello World\n");
}

int main() {
    Mykernel<<<1,32>>>();
    return 0;
}
```

Output:

Hello World

Hello World

.

32 times

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
__global__ void Mykernel() {
    printf("Hello World\n");
}
int main() {
    Mykernel<<<x,y>>>();
    return 0;
}
```

Output:

Hello World

Hello World

.

x * y times

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
#define N 10
__global__ void Mykernel() {
    printf("%d\n",N);
}
int main() {
    Mykernel<<<red>red>>>();
    return 0;
}
```

More GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
#define N 10
__global__ void Mykernel() {
    printf("%d\n",N);
}

int main() {
    Mykernel<<<x,y>>>();
    return 0;
}
```

Output:

10

10

.

x * y times

GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>

Const char *data = "Hello World";

__global__ void Mykernel() {
    printf("%s\n",data);
}

int main() {
    Mykernel<<<1,1>>>();
    return 0;
}
```

GPU Hello World program

```
#include <stdio.h>
#include <cuda.h>
Const char *data = "Hello World";
__global__ void Mykernel() {
    printf("%s\n",data);
}

int main() {
    Mykernel<<<1,1>>>();
    return 0;
}
```

Output:

ERROR

identifier "data" is undefined in
device code

Thank You