

Parallel Computing Workshop

Ms. Kranti Ingale

Prime Minister's Research Fellows (**PMRF**) Scholar

CSE, IIT Madras

Roadmap

- Introduction to GPU
- CUDA Program Flow and CPU-GPU Communication
- Thread organization (Grids, Blocks, Threads, 1D/2D)
 - WARP & Thread Divergence
- CUDA Memory Model
- CUDA Functions
- CUDA Thrust

Roadmap

- Introduction to GPU
- CUDA Program Flow and CPU-GPU Communication
- Thread organization (Grids, Blocks, Threads, 1D/2D)
 - WARP & Thread Divergence
- CUDA Memory Model
- CUDA Functions
- CUDA Thrust

Recap : Standard Variables

- Dimension of Grid (Number of blocks in the grid)
`gridDim.x , gridDim.y , gridDim.z`
- Dimension of block (Number of threads in a block)
`blockDim.x , blockDim.y , blockDim.z`

Recap : Standard Variables

- Block Index

`blockIdx.x , blockIdx.y,blockIdx.z`

- Thread Index

`threadIdx.x , threadIdx.y, threadIdx.z`

Recap : Example

- dim3 nblock(3,4,1) ;
- dim3 threadPerBlock(5,4,1) ;
- dkernel<<<nblock, threadPerBlock>>>();
- dkernel<<<(3,4,1), (5,4,1)>>>();

Recap : Example

- `dkernel<<<(3,4,1), (5,4,1)>>>();`
- `gridDim.x = 3 ; gridDim.y = 4 ; gridDim.z = 1`
- `blockDim.x = 5 ; blockDim.y = 4 ; blockDim.z = 1`
- `blockIdx.x = 0..2 ; blockIdx.y = 0..3 ; blockIdx.z = 0`
- `threadIdx.x = 0..4 ; threadIdx.y = 0..3 ; threadIdx.z = 0`

Matrix Multiplication

$$\text{Matrix 1} \begin{Bmatrix} 1 & 1 & 1 \\ 2 & 2 & 2 \\ 3 & 3 & 3 \end{Bmatrix} \quad \text{Matrix 2} \begin{Bmatrix} 1 & 1 & 1 \\ 2 & 2 & 2 \\ 3 & 3 & 3 \end{Bmatrix}$$

$$\text{Matrix 1} \begin{matrix} * \\ \text{Matrix 2} \end{matrix} \begin{Bmatrix} 1*1+1*2+1*3 & 1*1+1*2+1*3 & 1*1+1*2+1*3 \\ 2*1+2*2+2*3 & 2*1+2*2+2*3 & 2*1+2*2+2*3 \\ 3*1+3*2+3*3 & 3*1+3*2+3*3 & 3*1+3*2+3*3 \end{Bmatrix} =$$

$$\text{Matrix 1} \begin{matrix} * \\ \text{Matrix 2} \end{matrix} \begin{Bmatrix} 6 & 6 & 6 \\ 12 & 12 & 12 \\ 18 & 18 & 18 \end{Bmatrix}$$

Matrix Multiplication

```
void Mult(unsigned *matrix1,Unsigned *matrix2, unsigned *result, unsigned n)
{
for (unsigned i = 0; i < n; i++) {
    for (unsigned j = 0; j < n;j++)  {

        for (unsigned k = 0; k < n; k++) {
            result[i * n + j] += matrix[i * n + k] * matrix[k * n + j] ;
    } } } }
```

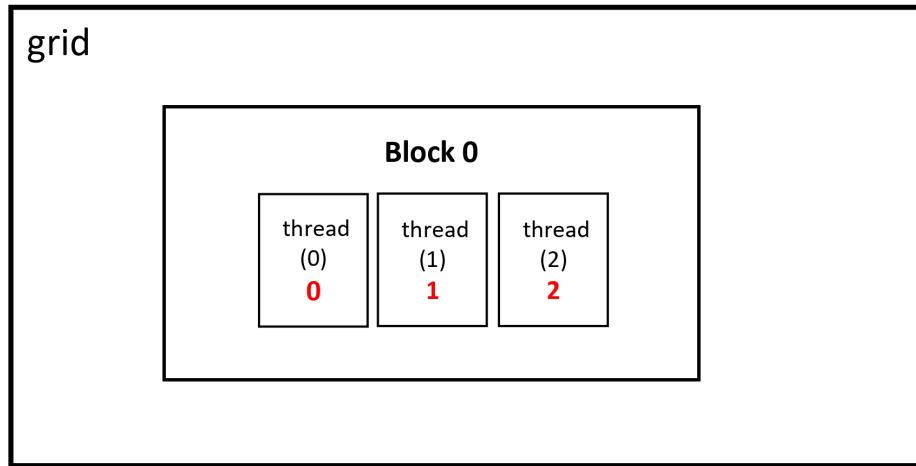
Matrix Multiplication

```
__global__ Mult(unsigned *matrix1,Unsigned *matrix2, unsigned *result, unsigned n)
{
    int id = blockDim.x * blockIdx.x + threadIdx.x;

    for (unsigned j = 0; j < n;j++) {
        for (unsigned k = 0; k < n; k++) {
            result[id * n + j] += matrix1[id * n + k] * matrix2[k * n + j];
        }
    }
}

Mult<<<1,n>>>(A,B,C,n);
```

Matrix Multiplication



`Mult<<<1,n>>>(A,B,C,n); // n=3`

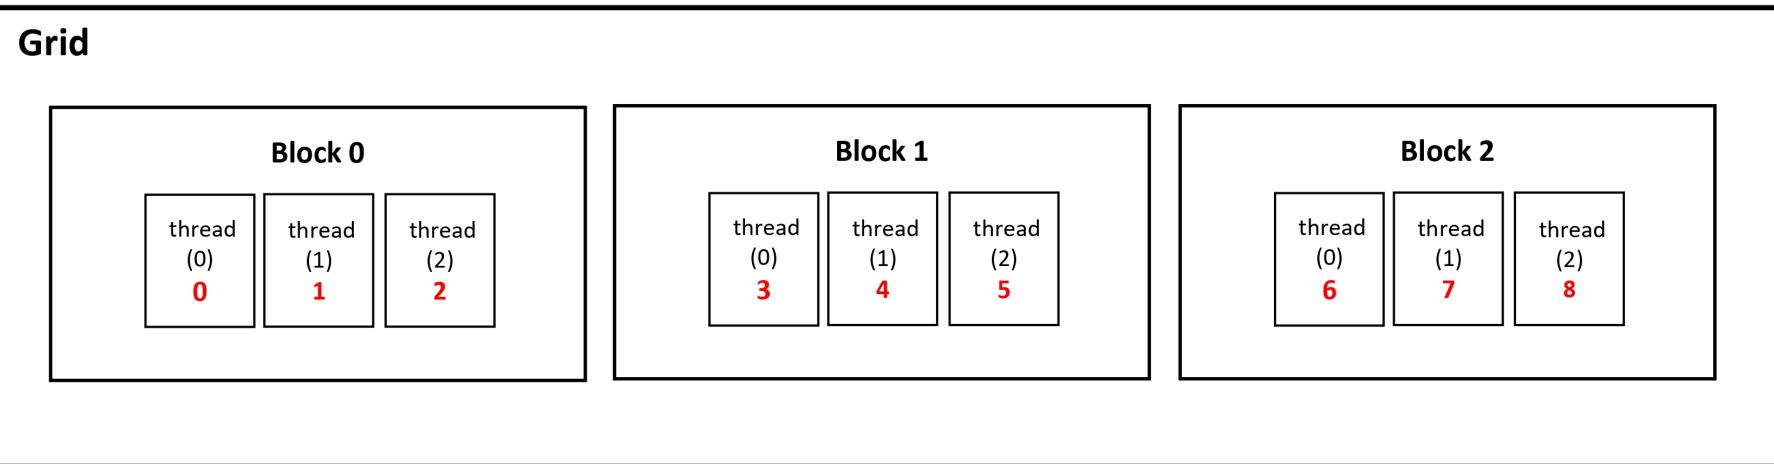
Matrix Multiplication

```
__global__ Mult(unsigned *matrix1,Unsigned *matrix2, unsigned *result, unsigned n)
{
    int id = blockDim.x * blockIdx.x + threadIdx.x;
    int i= id / n;
    int j= id % n;

    for (unsigned k = 0; k < n; k++) {
        result[id * n + j] += matrix1[id * n + k] * matrix2[k * n + j];
    } } }
```

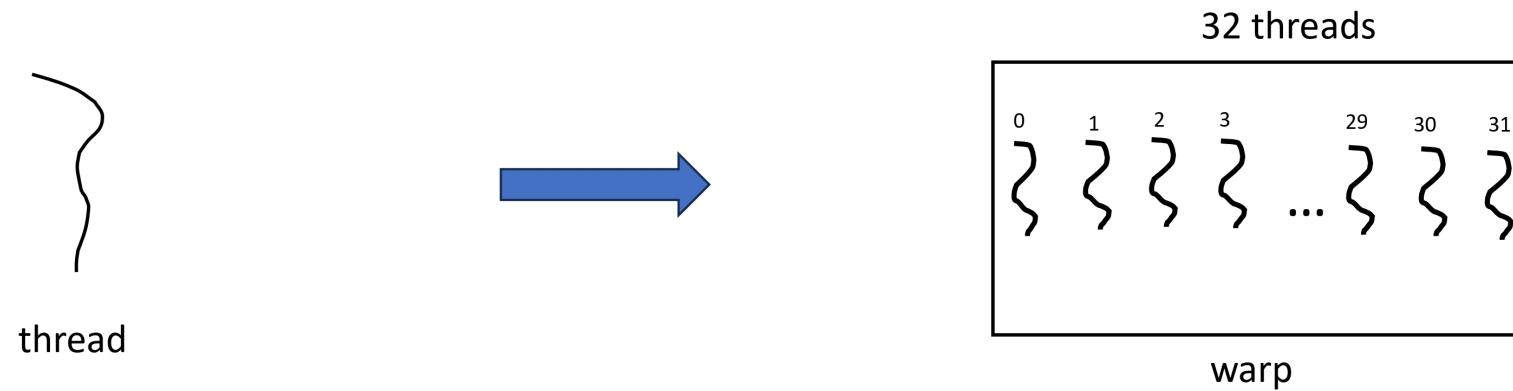
Mult<<<n,n>>>(A,B,C,n);

Matrix Multiplication

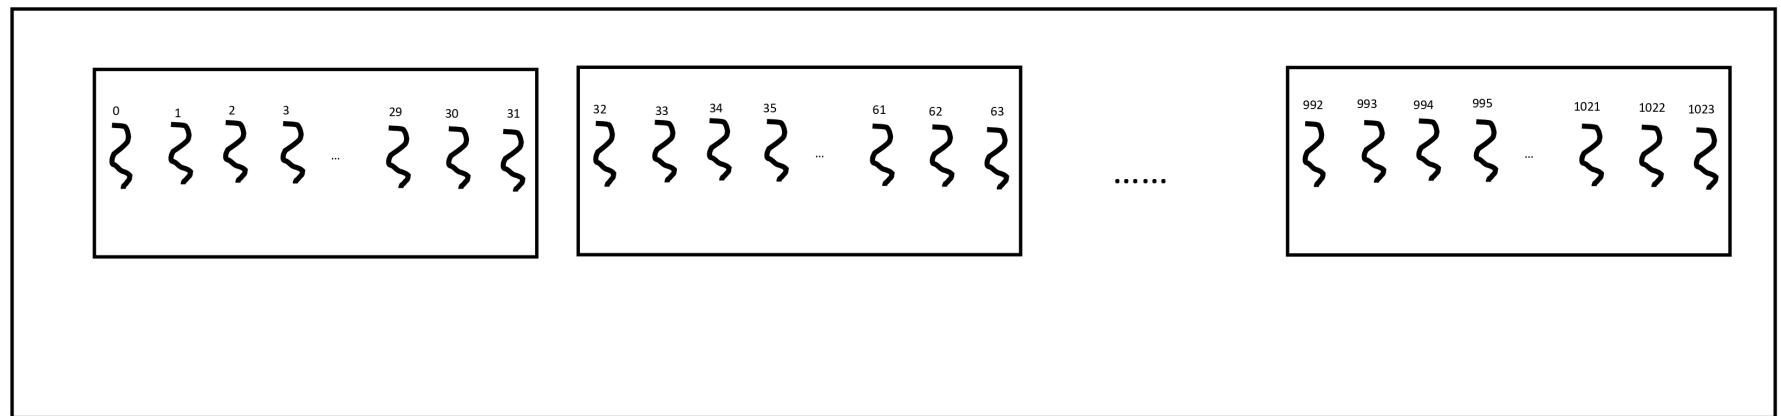


Mult<<<n,n>>>(A,B,C,n); // n=3

COMPUTATIONAL HIERARCHY

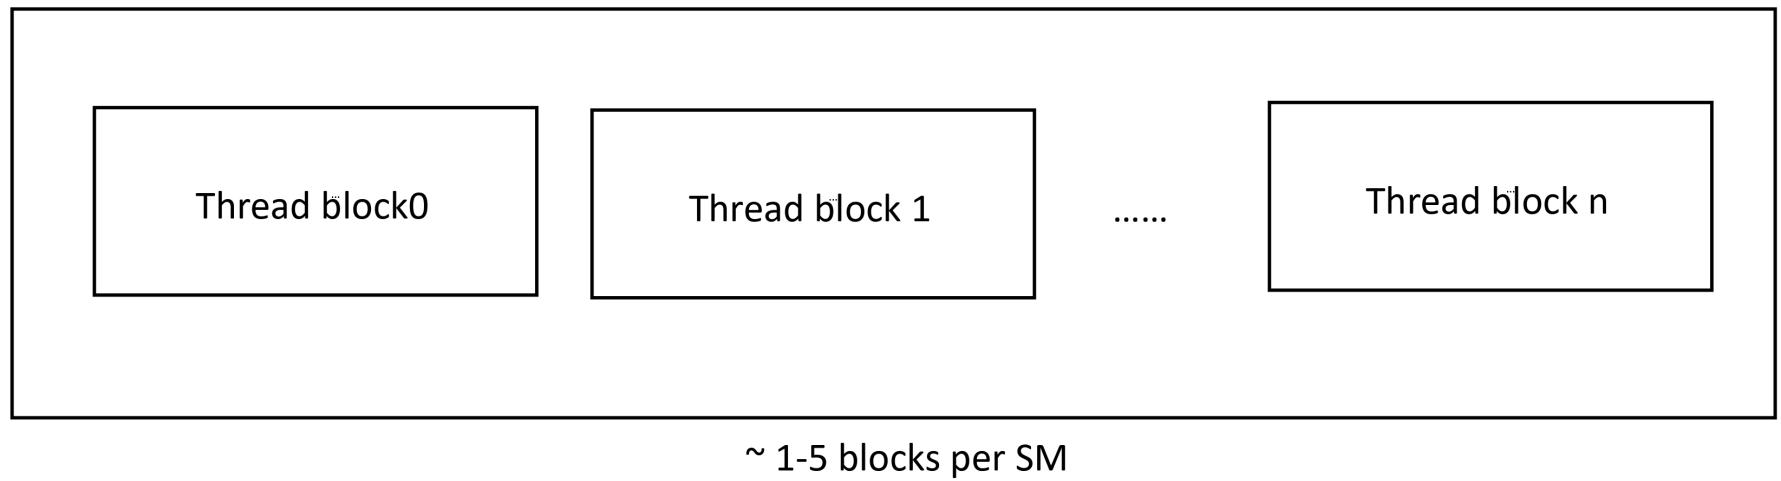


COMPUTATIONAL HIERARCHY

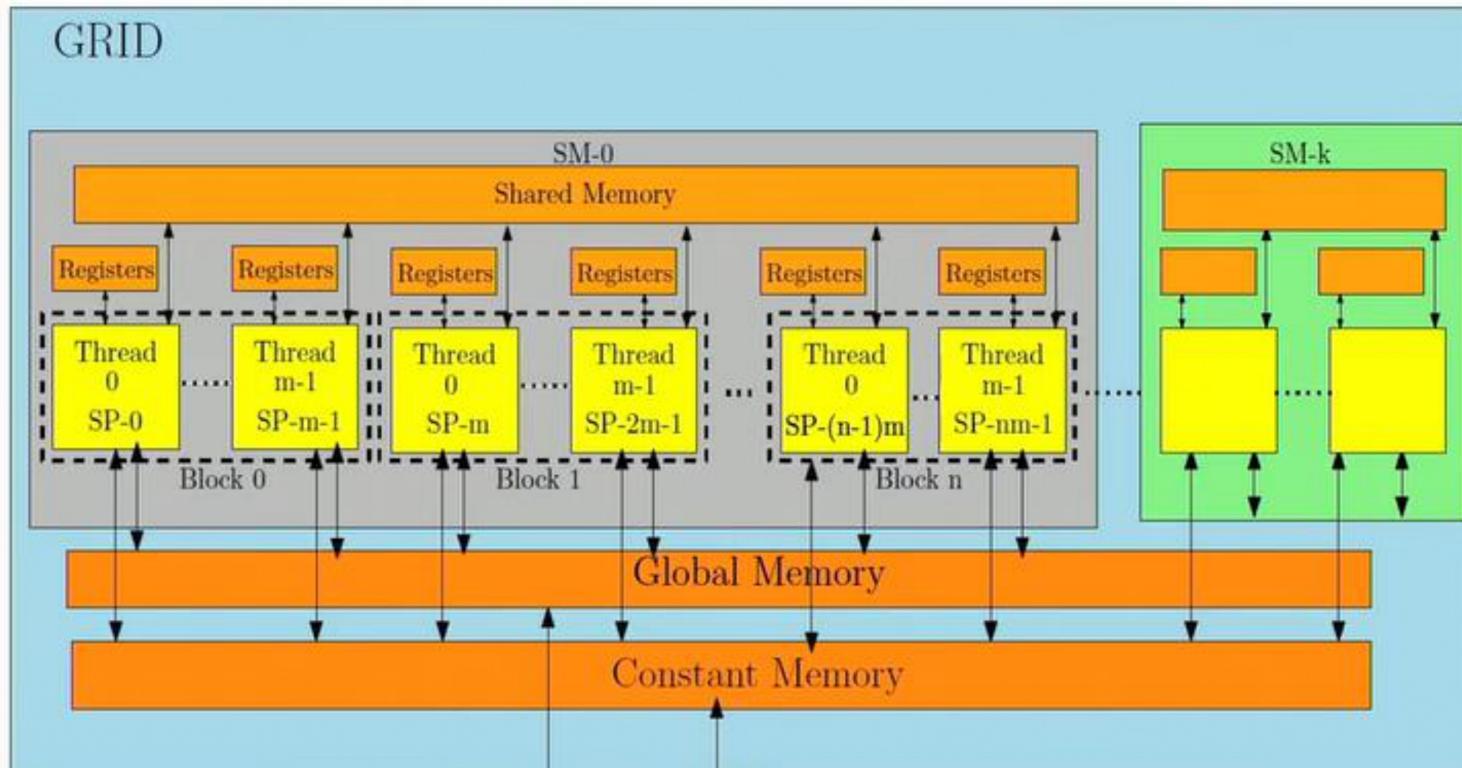


1 Thread Block = 1024 threads

COMPUTATIONAL HIERARCHY



Thread mapping to Hardware



$\sim 1\text{-}5$ blocks per SM

Example-1

- Grid < 6,2,1> Block = <5,4,1>
- Number of SM's = 6
- SP's per SM = 40

Example-1

- Grid < 6,2,1> Block = <5,4,1>
- Number of SM's = 6
- SP's per SM = 40
- Number of thread block = $6 * 2 = 12$

Example-1

- Grid < 6,2,1> Block = <5,4,1>
- Number of SM's = 6
- SP's per SM = 40
- Number of thread block = $6 * 2 = 12$
- Threads per thread_block = $5 * 4 = 20$
- Two thread block are mapped to one SM

Thread mapping to Hardware

